BUYSURE AI: An Intelligent Shopping Assistant for Trustworthy Product Evaluation and Personalized Purchase Decision Support

Srikowsalya R, Sujith S, Lingesh S
Department of Computer Science and Engineering
Vels Institute of Science, Technology and Advanced Studies (VISTAS), Chennai, India
Srikowsalya R: srikowsalya06@gmail.com Sujith S: sxjithh@gmail.com Lingesh S: lingeshs256@gmail.com
Project Guide: Janani P, Assistant Professor

***Abstract*—Online shopping provides consumers with an enormous selection of products, but the size of the marketplace also increases the difficulty of making a reliable purchase decision. Product pages may contain hundreds of reviews, repeated or low-information comments, conflicting opinions, changing prices, similar competing products, and seller information that is difficult to interpret quickly. BUYSURE AI is proposed as an AI-based shopping assistant that converts these scattered signals into a structured and understandable decision-support report. The system combines product information, customer reviews, ratings, prices, product features, comparable alternatives, and user-defined requirements. A Natural Language Processing (NLP) layer is used to preprocess and interpret review text, estimate sentiment and aspect-level opinions, and identify suspicious or potentially misleading review patterns. A Machine Learning (ML) layer can combine textual, rating, metadata, and behavioral indicators to estimate review reliability. A comparison engine normalizes product attributes and evaluates competing products against a user's budget and priorities. The final decision-support layer summarizes the evidence using an interpretable recommendation such as *Buy*, *Wait*, or *Avoid*, together with the main positive factors, negative factors, risks, and trade-offs. The architecture also provides a path for future capabilities such as screenshot and image understanding, price tracking, seller-risk analysis, and automatic discovery of better alternatives. The central objective is not to replace human judgment, but to reduce information overload and improve the transparency, consistency, and usefulness of online purchase decisions.**

***Index Terms*—Artificial intelligence, e-commerce, fake review detection, natural language processing, sentiment analysis, machine learning, recommender systems, product comparison, decision support, personalized shopping.**

I. Introduction

Online commerce has changed the way consumers discover, compare, and purchase products. A buyer can now evaluate several brands, models, prices, ratings, specifications, and customer opinions without visiting a physical store. This convenience, however, creates a new decision-making problem: the amount of information available can be much larger than the amount of information a user can reasonably inspect before purchasing. The BUYSURE AI project is motivated by this gap between information availability and information usability. The supplied project abstract specifically identifies fake reviews, misleading product information, fluctuating prices, unreliable sellers, and the large number of similar products as major sources of uncertainty.

Traditional shopping decisions often depend on an average star rating and a quick reading of a few reviews. Such a strategy can be misleading. A high average rating does not explain whether the positive comments are detailed and relevant, whether negative experiences concern an important feature, or whether the review set contains suspicious patterns. Likewise, the lowest price is not necessarily the best value if the product has weak specifications, poor seller reliability, limited warranty support, or repeated complaints about a critical feature. These observations motivate a system that treats a product page as a collection of heterogeneous evidence rather than a single rating number.

The supplied BUYSURE AI abstract defines the intended system as an AI-based shopping assistant using ML and NLP to analyze product details, reviews, ratings, prices, features, and comparable products. It also specifies that the system should identify suspicious or potentially misleading reviews, extract positive and negative aspects from genuine feedback, compare products according to price, features, quality, ratings, and user requirements, and provide an understandable recommendation such as *Buy*, *Wait*, or *Avoid*. User budget, preferences, and specific requirements are included as personalization inputs. [1]

This paper converts that project concept into a detailed conference-paper architecture. The proposed design separates data acquisition, preprocessing, review intelligence, product normalization, comparison, personalization, decision generation, and explanation. This separation is important because each stage has a different responsibility and can be evaluated independently. For example, a review classifier should be evaluated as a review-analysis component, while the final recommendation should be evaluated as a decision-support component. A correct classification of a suspicious review does not automatically guarantee that the final purchase recommendation is correct.

The design is intentionally explainable. Instead of returning only a label, the system is designed to expose the factors behind the recommendation: price fit, requirement match, product

strengths, product weaknesses, review confidence, seller risk indicators, and comparison with alternatives. This approach is consistent with research showing that online reviews contain useful product-level and feature-level information that can support recommendation systems when appropriate text-mining and review-processing methods are applied. [6]

A. Motivation and Problem Statement

The main problem addressed by BUYSURE AI is the difficulty of converting a large and noisy collection of e-commerce information into a trustworthy purchasing summary. Let a product be represented by

$$P = \{D, F, R, T, S, C\}, \tag{1}$$

where D is descriptive product information, F is structured features, R is the review set, T represents price and temporal information, S represents seller-related information, and C represents comparable products. The user profile is represented as

$$U = \{B, W, Q, P_r\}, \tag{2}$$

where B is budget, W is feature-weight preference, Q is explicit requirements, and P_r is other user priorities.

The system must transform (P, U) into an evidence-based decision report E rather than merely a binary recommendation:

$$E = f(D, F, R, T, S, C, U), \tag{3}$$

where E contains a recommendation, supporting evidence, risks, and comparison information. In the current project description, the exact production dataset, model architecture, trained parameters, and deployment environment are not specified. Therefore, the technical details in this paper describe the proposed system architecture and processing methodology rather than claiming completed benchmark performance.

B. Objectives

The major objectives are: (1) collect and organize product information from an available source; (2) preprocess and analyze customer reviews; (3) detect suspicious review patterns using ML/NLP methods; (4) extract positive and negative product aspects from review text; (5) compare multiple products using normalized features and user priorities; (6) personalize the evaluation according to budget and requirements; (7) produce an interpretable Buy/Wait/Avoid decision; and (8) preserve a modular architecture so that future image analysis, price tracking, seller-risk analysis, and alternative discovery can be integrated without redesigning the complete system.

C. Paper Organization

The remainder of this paper first reviews related work, then defines the proposed architecture, data pipeline, review-analysis engine, comparison and recommendation logic, interface workflow, implementation plan, evaluation methodology, limitations, security considerations, and future enhancements. The final section summarizes the contribution and practical value of BUYSURE AI.

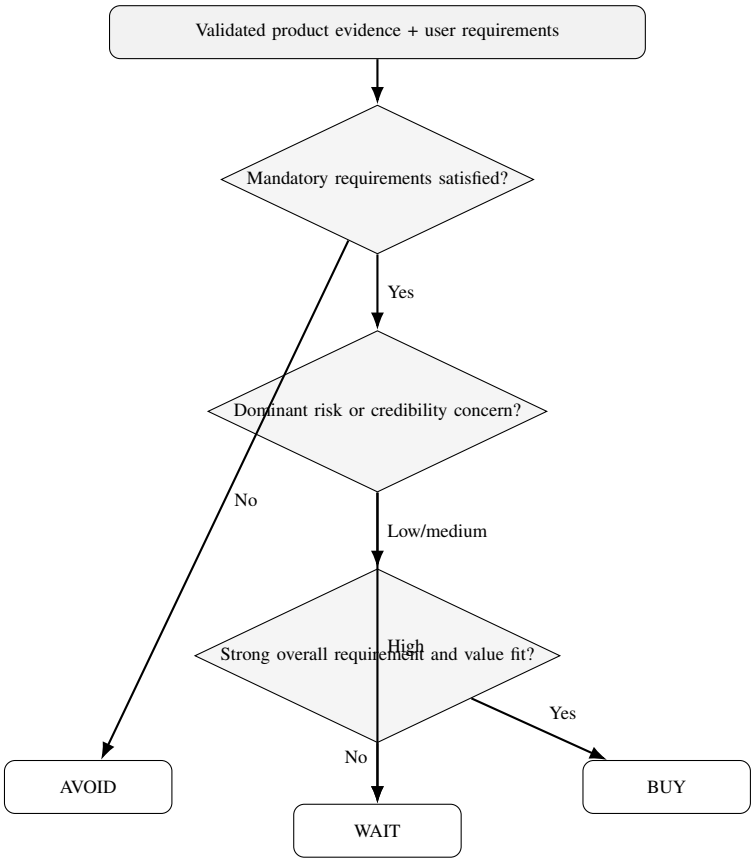


Fig. 1: Proposed explainable decision flow for Buy/Wait/Avoid classification. Exact thresholds are implementation parameters to be evaluated rather than assumed.

II. Related Work and Research Background

Fake-review detection has become an important research problem because consumers frequently use online reviews as evidence before purchasing. A major IEEE survey describes fake-review detection as a problem involving linguistic, behavioral, and statistical signals and compares traditional machine-learning and deep-learning approaches. The study also reports that transformer-based representations can provide strong baselines for deceptive-review classification, illustrating the value of contextual language representations. [2]

More recent work continues to examine combinations of NLP and ML for e-commerce review credibility. An IEEE conference study published in 2025 describes a framework that combines textual features such as sentiment polarity and n-gram information with reviewer behavior and temporal patterns. The study reports the use of classifiers including KNN, logistic regression, and SVM, illustrating a useful design principle for BUYSURE AI: review credibility can be treated as a multi-signal problem rather than a purely textual classification problem. [4]

Deep-learning surveys further demonstrate the growing use of neural methods for deceptive-review detection. Such approaches can learn contextual relationships that are difficult to represent using simple keyword counts. However, a complex model does not automatically guarantee trustworthy decisions; class imbalance, domain changes, generated text, adversarial behavior, and differences between datasets remain important challenges. [3]

A hybrid review-detection study has also emphasized the importance of review representation. Attention-based recurrent architectures can focus on informative words or phrases, while imbalance-aware training can help when genuine and deceptive examples are unevenly distributed. These ideas support the proposed architecture's separation between preprocessing, feature representation, classification, and explanation. [5]

Recommendation research provides another foundation. Reviews are not merely ratings; they contain fine-grained descriptions of specific product attributes, strengths, and weaknesses. Research and tutorials on review-enhanced recommendation describe techniques such as sentiment analysis, topic modeling, distributed representations, and recommender algorithms that directly exploit review information. [6] More broadly, multi-criteria review-based recommender research argues that a single overall rating cannot capture fine-grained user preferences, while explainable recommender research emphasizes the importance of understandable reasons for recommendations. [7], [8] More broadly, multi-criteria review-based recommender research argues that a single overall rating cannot capture fine-grained user preferences, while explainable recommender research emphasizes the importance of providing understandable reasons for recommendations. [7], [8]

A. Research Gap and Position of BUYSURE AI

Existing research commonly focuses on one subproblem, such as fake-review classification or recommendation. BUYSURE AI combines several decision-support functions in one user-facing workflow. The proposed system does not claim to invent a new fake-review algorithm. Instead, its contribution is the integration of review credibility, sentiment and aspect extraction, product comparison, user requirements, and an explainable Buy/Wait/Avoid output.

Table I summarizes the relationship between the research areas and the proposed project.

TABLE I: Research Areas and Their Role in BUYSURE AI

Area	Established direction	Role in proposed system
Fake-review detection	ML, deep learning, linguistic and behavioral features	Estimate review credibility and flag suspicious evidence.
Sentiment analysis	Polarity and opinion mining	Identify positive, negative, and mixed feedback.
Aspect extraction	Fine-grained product-feature opinions	Convert free text into feature-level strengths and weaknesses.
Recommendation	Ratings, content, reviews, user preferences	Rank products according to user priorities.
Explainable AI	Evidence and feature attribution	Show why a product receives its recommendation.

III. Proposed BUYSURE AI Architecture

A. System Overview

The proposed architecture is organized into seven logical layers: user interaction, product-data acquisition, preprocessing and normalization, review intelligence, product comparison, decision and explanation, and optional future services. The architecture is designed so that raw product information is not sent directly to a recommendation model. Instead, data are validated, normalized, analyzed, and converted into intermediate evidence before the final decision is generated.

B. Layer 1: User Interaction

The user specifies the information that makes a recommendation meaningful. At minimum, the interface should accept a budget range and one or more requirements. Requirements can be represented as mandatory constraints, preferred features, or soft preferences. For example, a user may require a product below a certain price, prefer a strong camera, and give battery life a higher priority than appearance. Separating mandatory requirements from preferences prevents the system from treating every feature as equally important.

C. Layer 2: Product Data Acquisition

The acquisition layer collects the available product title, category, brand, price, rating, feature list, seller information, and review content. Because online product pages are heterogeneous, the architecture uses a normalized internal representation. Each product record should retain its original values as well as a normalized value when possible. This allows the recommendation engine to reason over comparable units while preserving the original information for display.

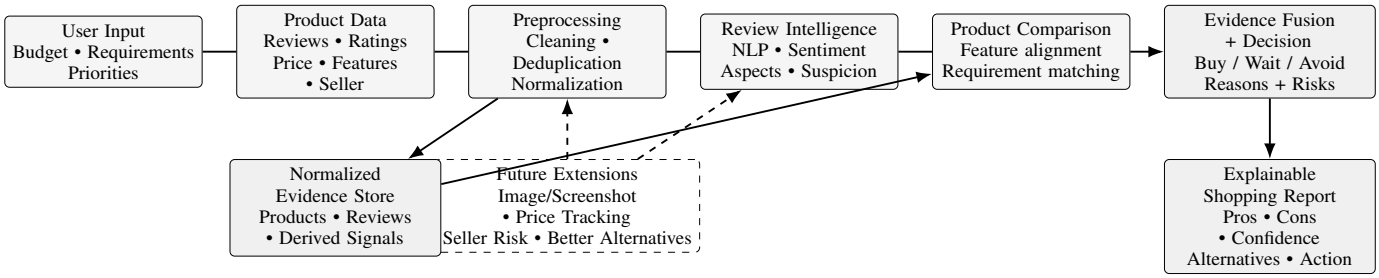


Fig. 2: Detailed proposed BUYSURE AI architecture. The solid path represents the core decision-support pipeline; the dashed extension layer represents future capabilities identified in the project abstract.

D. Layer 3: Preprocessing and Normalization

Text preprocessing removes irrelevant formatting and prepares reviews for NLP. Typical operations include Unicode normalization, whitespace cleanup, sentence segmentation, tokenization, and duplicate detection. Depending on the selected model, stop-word removal and stemming may be used, although transformer-based systems generally benefit from model-specific tokenization rather than aggressive linguistic reduction.

Numerical normalization is equally important. A price in rupees, a storage capacity in gigabytes, and a battery capacity in milliampere-hours cannot be combined directly. Each field is converted into a common internal representation, while the original unit is retained for display. Categorical values such as warranty type or operating system can be mapped to standardized labels. Missing values are preserved as missing rather than silently assigned a favorable value.

E. Layer 4: Review Intelligence

The review engine operates on individual reviews and aggregates evidence at product level. It contains three conceptual functions: credibility analysis, sentiment analysis, and aspect extraction. Credibility analysis identifies suspicious patterns; sentiment analysis estimates the polarity or tone of feedback; and aspect extraction associates opinions with product attributes such as battery, camera, display, delivery, durability, or performance. The exact model can be replaced independently of the rest of the architecture.

F. Layer 5: Product Comparison

The comparison engine aligns products on common features and evaluates them against user requirements. A product should not receive a high score merely because it has more features. Features are useful only when they contribute to the user's stated objective. Thus, the comparison layer supports both hard constraints and weighted preferences.

G. Layer 6: Evidence Fusion and Decision Support

The decision layer combines normalized product quality signals, review evidence, price fit, requirement match, and comparison results. The output is not intended to be an absolute truth statement. It is a decision-support classification that summarizes available evidence and communicates uncertainty. The explanation generator attaches reasons to the recommendation so that users can inspect the basis of the result.

H. Layer 7: Future Extensions

The supplied abstract identifies future possibilities including product image and screenshot analysis, price tracking, seller-risk detection, and automatic identification of better alternatives. These capabilities are represented as extensions rather than completed modules. Their placement at the edge of the architecture allows them to contribute additional evidence without changing the fundamental decision pipeline. [1]

I. Architecture Rationale and Data Provenance

The architecture deliberately separates raw inputs from derived evidence. A raw review should remain available for traceability, while the processed sentiment, extracted aspects, and credibility indicators are stored as derived fields. This separation makes it possible to audit why a recommendation changed when a new review or price observation is introduced. It also prevents a later model update from destroying the original evidence used to produce an earlier decision.

A second design principle is modular replacement. The review model can be upgraded without changing the product-comparison layer, and a new comparison strategy can be evaluated without changing the user interface. Each module should expose a small, documented input/output contract. For example, the review engine can return a review identifier, credibility indicator, sentiment, aspect list, and confidence values. The decision engine can consume these fields without knowing the internal neural architecture that produced them.

The third principle is provenance. Every derived signal should be associated with its source record and analysis version. If a product receives an Avoid recommendation because several negative reviews were influential, the system should be able to identify which reviews contributed to that conclusion and whether those reviews were considered credible. This is especially important when the system is used as a decision-support tool rather than a simple product-ranking page.

TABLE II: Proposed Product Data Representation

Field group	Example contents	Purpose
Identity	Product ID, title, brand, category	Unique product reference.
Commercial	Current price, list price, availability	Price-fit and timing analysis.
Features	Capacity, display, battery, processor, etc.	Requirement matching.
Ratings	Average rating, rating count	Aggregate reputation signal.
Reviews	Text, rating, date, metadata	NLP and credibility analysis.
Seller	Seller name, rating, policy indicators	Seller-risk assessment when data exist.
User profile	Budget, priorities, mandatory needs	Personalized scoring.
Derived evidence	Sentiment, aspects, suspicion score	Decision-support inputs.

IV. Data Model, Preprocessing, and Review Intelligence

A. Unified Product Representation

A practical implementation requires a consistent internal representation because product pages frequently use different names and formats for equivalent attributes. The proposed schema contains identity fields, commercial fields, feature fields, review fields, and user-specific evaluation fields.

B. Review Preprocessing Pipeline

The review pipeline is intentionally staged because a classifier trained on unprocessed text may learn irrelevant formatting patterns. First, the review is validated and assigned an internal identifier. Second, text is cleaned while preserving information that may be useful for linguistic analysis. Third, the text is segmented into sentences and tokens. Fourth, representations are created for sentiment, aspect extraction, and suspicious-review analysis. Finally, the outputs are stored as evidence rather than immediately converted into a final product decision.

C. Suspicious Review Detection

The supplied project description requires the system to identify suspicious or potentially misleading reviews. This should be treated as a risk-estimation problem rather than an accusation that a specific reviewer is fraudulent. A review can be flagged because its linguistic or behavioral signals are unusual, but a flag does not prove deception. The user interface should therefore use terms such as *suspicious*, *low-confidence*, or *credibility concern* rather than declaring a person dishonest.

A proposed review feature vector can be represented as

$$X_r = [L, S, N, T, D, B, M], \tag{4}$$

where L represents linguistic characteristics, S sentiment features, N n-gram or semantic features, T temporal behavior, D duplication/similarity evidence, B reviewer behavior when available, and M other metadata. The exact feature set depends on the data source.

A model can estimate a review credibility probability

$$p_{cred}(r) = P(\text{genuine} \mid X_r). \tag{5}$$

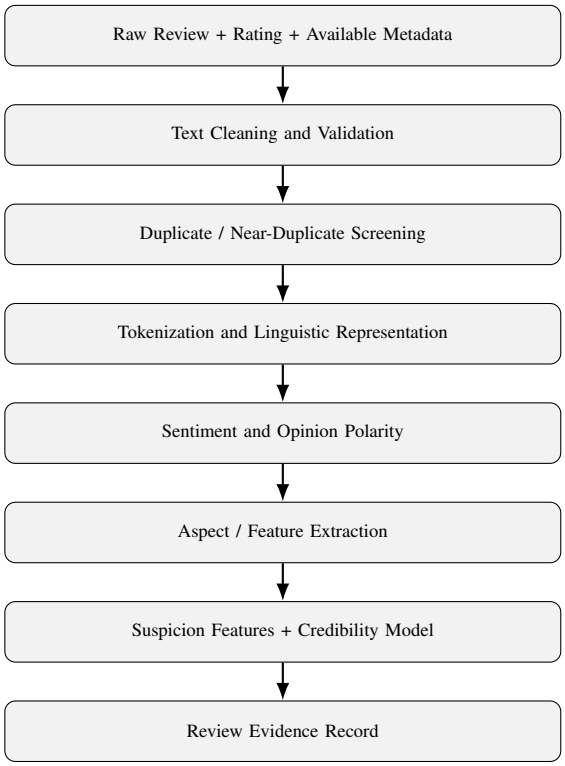


Fig. 3: Detailed review-analysis pipeline proposed for BUYSURE AI.

The output should be calibrated before being interpreted as a probability. If the project uses a simpler classifier, the system can instead expose a risk score with an explicit label such as low, medium, or high concern.

Recent literature supports combining textual and non-textual information in fake-review detection rather than relying on sentiment alone. [2], [4]

D. Sentiment and Aspect Extraction

Sentiment analysis identifies whether a review expresses positive, negative, neutral, or mixed opinion. However, product-level sentiment is more useful when linked to specific aspects. For example, a review may be positive about camera quality but negative about battery life. The proposed system therefore represents a review as a set of aspect-opinion pairs:

$$A_r = \{(a_1, o_1, s_1), \ldots, (a_k, o_k, s_k)\}, \tag{6}$$

where a_i is an aspect, o_i is the associated opinion phrase, and s_i is a sentiment or confidence value.

The product-level aspect summary can then aggregate evidence across reviews while weighting reviews by credibility. A simple proposed formulation is

$$S_a = \frac{\sum_{r \in R_a} w_r s_{r,a}}{\sum_{r \in R_a} w_r}, \tag{7}$$

where R_a is the set of reviews mentioning aspect a, w_r is a credibility weight, and $s_{r,a}$ is the sentiment contribution. This formulation is a proposed design, not a reported experimental result.

E. Handling Conflicting Reviews

Conflicting feedback should not be hidden. Suppose ten reviews praise battery life while three credible reviews report rapid battery drain under heavy use. A simple average may hide the context. The system can instead show a summary such as: *Battery feedback is mostly positive, but a smaller set of credible reviews reports heavy-use drain.* This explanation is more informative than a single sentiment score because it preserves disagreement.

F. Review Evidence Aggregation

For each product, the review engine should retain: number of reviews analyzed, fraction of reviews considered suspicious, sentiment distribution, major positive aspects, major negative aspects, and confidence in each aspect. This intermediate representation becomes the input to the comparison engine. The architecture therefore avoids a common failure mode in which a complex NLP model produces a score that cannot be traced back to the user-facing recommendation.

V. Product Comparison and Decision Engine

A. Requirement Matching

The user profile is converted into a set of constraints and weighted preferences. A mandatory requirement is evaluated using a pass/fail rule. A preference is represented using a weight w_i such that

$$\sum_{i=1}^{n} w_i = 1, \qquad w_i \geq 0. \tag{8}$$

For product p, the normalized value of feature i is $x_{p,i}$. A requirement-match score can be expressed as

$$M(p) = \sum_{i=1}^{n} w_i x_{p,i}. \tag{9}$$

The normalization function depends on whether a higher value is desirable, a lower value is desirable, or the attribute is categorical.

B. Price Fit

Price should be interpreted relative to the user's budget. A product inside the budget is not automatically superior, but exceeding a hard budget can make a product unsuitable. A proposed budget-fit function is

$$B(p) = \begin{cases} 1, & price(p) \leq B_{max}, \\ 0, & price(p) > B_{max}, \end{cases} \tag{10}$$

for a strict budget. For flexible budgets, a continuous penalty can be used instead. The interface should clearly state whether the user's budget was treated as a hard constraint or a preference.

TABLE III: Illustrative Decision Evidence Matrix

Evidence	Positive indication	Concern indication
Budget fit	Within stated budget	Above hard limit
Requirement match	Mandatory needs satisfied	One or more mandatory needs missing
Review credibility	Mostly credible evidence	High suspicious-review proportion
Aspect sentiment	Strong positive feedback on priority features	Repeated negative feedback on priority features
Rating consistency	Rating supported by detailed reviews	Rating inconsistent with review evidence
Price trend	Stable/favorable current value	Recent increase or uncertain value when tracking exists
Seller evidence	Strong seller indicators	Risk indicators or insufficient evidence
Alternatives	Product ranks strongly against comparable items	Better-matching alternative exists

C. Evidence-Fusion Score

A transparent product score can combine requirement match, review evidence, rating evidence, value-for-money, and risk. One proposed form is

$$Score(p) = w_m M(p) + w_r R(p) + w_v V(p) + w_q Q(p) - w_k K(p), \tag{11}$$

where M is requirement match, R is review evidence, V is value-for-money, Q is overall quality evidence, and K is a risk penalty. The weights are configurable and should be documented during evaluation.

The equation is deliberately interpretable. If a user prioritizes battery life, the system should increase the contribution of the battery-related feature rather than silently changing the entire recommendation model. If seller risk becomes significant, the risk term can lower the decision score while the explanation identifies the reason.

D. Buy, Wait, and Avoid Logic

The final label is generated after the evidence has been combined. *Buy* should indicate that the product satisfies the important user constraints and that the available evidence contains no dominant concern. *Wait* should indicate that the product is potentially suitable but that timing, price uncertainty, insufficient evidence, or a meaningful unresolved trade-off weakens the immediate decision. *Avoid* should indicate a strong mismatch, a major requirement failure, or a substantial risk signal.

These labels are not medical, legal, or financial advice and should not be presented as guarantees. They summarize the system's available evidence. The user should always be able to inspect the reasons behind the label.

VI. Personalization and Explainable Recommendation

A. Why Personalization is Required

A universal ranking is inadequate because users have different constraints. Two users can inspect the same smartphone and reach different correct decisions: one may prioritize camera performance, another battery endurance, and a third low price. BUYSURE AI therefore treats the user's priorities as first-class inputs.

B. Preference Weighting

The interface can provide simple controls such as *most important*, *important*, and *less important*. These qualitative selections can be mapped to normalized numerical weights. The system should display the selected priorities before generating the recommendation so that the user knows what the model is optimizing.

C. Constraint Satisfaction Before Ranking

Hard constraints should be evaluated before soft preferences. If the user requires a product below a maximum budget, a product that exceeds that limit should not outrank an in-budget product merely because it has additional features. Similarly, if a user requires a particular compatibility feature, the absence of that feature should be treated as a decisive mismatch rather than a small negative score.

D. Recommendation Explanation

The explanation module translates internal evidence into short, human-readable reasons. A recommendation report can contain the following blocks:

1) **Decision:** Buy, Wait, or Avoid.
2) **Why:** the three to five strongest evidence points.
3) **Best features:** aspects with strong positive evidence.
4) **Concerns:** negative aspects, suspicious-review signals, or missing evidence.
5) **Budget view:** current price relative to the user's limit.
6) **Alternative view:** one or more products that better match the user's priorities when available.
7) **Confidence:** an indication of how much reliable evidence is available, without presenting confidence as certainty.

E. Example Decision Narrative

For illustration, suppose a user enters a budget and identifies battery life and camera quality as high-priority requirements. The system compares candidate products, finds that Product A has strong battery-related review evidence and acceptable camera feedback, while Product B has a lower price but repeated credible complaints about battery performance. A useful explanation would not simply state that Product A has a higher score. It would explain that Product A better satisfies the user's two priority features and has stronger supporting review evidence, while Product B's price advantage is offset by the recurring battery concern.

This narrative style is important because the project objective is informed decision making, not score maximization. The explanation should allow a user to disagree with the recommendation while still understanding how it was produced.

TABLE IV: Recommended User-Facing Report Sections

Section	Information shown
Decision card	Buy / Wait / Avoid, short rationale, evidence confidence.
Budget summary	Current price, budget fit, notable price concern if available.
Requirement match	Mandatory requirements and priority-feature match.
Review health	Review volume, suspicious-review indicators, sentiment balance.
Pros	Strongly supported positive aspects.
Cons	Repeated or important negative aspects.
Comparison	Relative strengths and weaknesses against selected alternatives.
Action	Suggested next step, such as compare, wait for price, or proceed.

F. Confidence and Missing Evidence

The system should distinguish between negative evidence and missing evidence. If seller information is unavailable, the report should say that seller risk could not be fully assessed rather than assigning a high risk automatically. Similarly, a product with very few reviews should not be treated as equivalent to a product with hundreds of detailed reviews. Evidence volume and evidence quality can therefore be included as metadata in the explanation.

VII. End-to-End User Workflow

The complete workflow begins when a user submits a product or product set for analysis. The system validates the input, retrieves or receives product information, normalizes the records, processes reviews, computes product-level evidence, applies the user's constraints and priorities, compares alternatives, and generates the final report.

A. User-Facing Output Structure

The final interface should prioritize readability. A recommended layout is a decision card at the top, followed by a compact score/evidence summary, major positive aspects, major concerns, product comparison, and a detailed review section. The system should avoid overwhelming the user with raw model probabilities. Instead, technical evidence can be expandable while the default view remains understandable.

VIII. Implementation Framework

The project can be implemented as a modular software stack. The exact programming language and framework are not specified in the supplied abstract, so the following is an implementation framework rather than a claim about the current codebase.

A. Application Layer

A web-based interface can provide the product input form, user preferences, comparison view, and final decision report. The interface communicates with a backend service through structured requests. A modular API boundary makes it possible to replace the NLP model without redesigning the user interface.

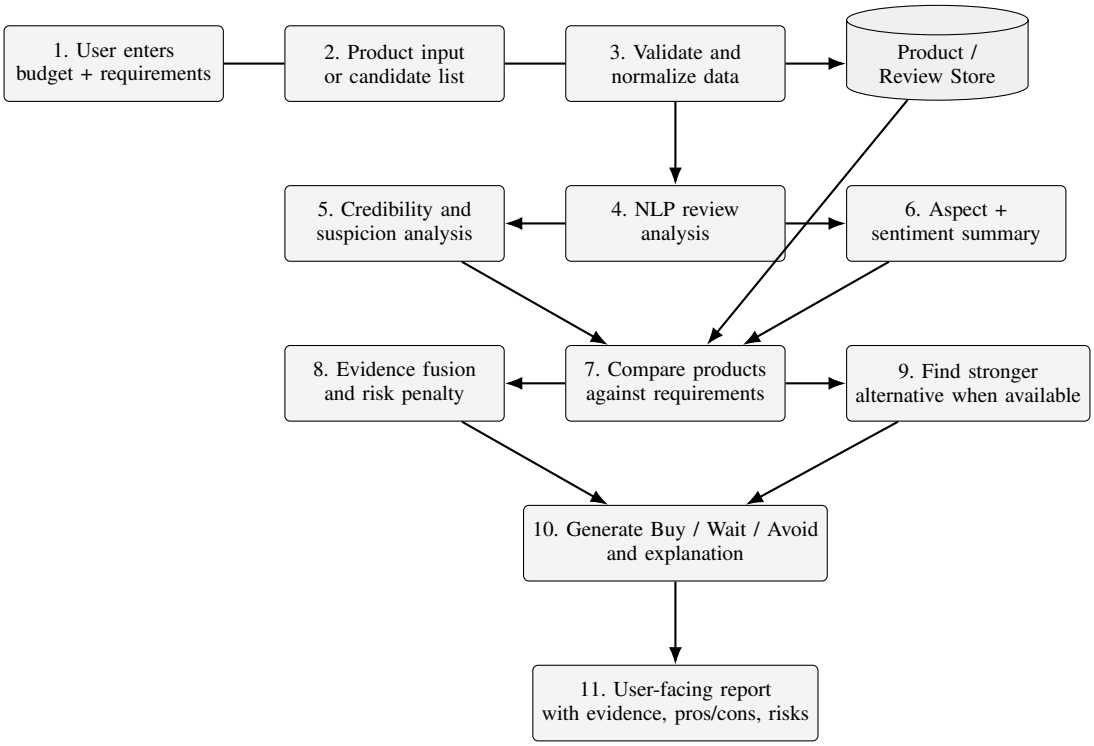


Fig. 4: End-to-end BUYSURE AI workflow from user requirements to an explainable purchase decision.

B. Data Layer

A document-oriented store can be used for product records and reviews because different product categories have different feature sets. A relational representation is also possible when the schema is sufficiently stable. The important architectural principle is to maintain raw input, normalized values, derived evidence, and recommendation output as distinguishable fields.

C. NLP and ML Layer

The NLP layer can contain separate components for text preprocessing, sentiment analysis, aspect extraction, and suspicious-review classification. Candidate model families include classical ML classifiers for engineered features and transformer-based encoders for contextual text representations. Model selection should depend on dataset size, compute resources, latency requirements, and the availability of labeled examples.

D. Decision Layer

The decision engine should not directly expose raw classifier probabilities as purchase advice. Instead, it should combine multiple evidence channels through documented rules or a calibrated scoring layer. This is the point at which user constraints are applied and where explanations are generated.

E. Suggested Software Modules

Table V provides a practical modular decomposition for development.

TABLE V: Suggested Software Modules

Module	Responsibility
Input Manager	Accept product URLs/data and user requirements.
Product Parser	Extract and standardize product attributes.
Review Processor	Clean, tokenize, and structure review text.
Review Credibility Engine	Generate suspiciousness/credibility evidence.
Sentiment Engine	Estimate review and aspect sentiment.
Feature/Aspect Engine	Map review opinions to product attributes.
Comparison Engine	Normalize and rank candidate products.
Decision Engine	Apply constraints, weights, and evidence fusion.
Explanation Engine	Convert evidence into human-readable reasons.
Report/UI Layer	Present decision, comparison, pros, cons, and risks.

F. Security and Data Protection

If user profiles or shopping histories are stored, the implementation should minimize personally identifiable information. User preferences should be stored only when necessary. Authentication, authorization, encrypted transport, input validation, and secure storage should be considered for any deployed version. Product scraping or external API access must also respect the terms and policies of the relevant source platform.

G. Failure Handling

A production implementation should define behavior for missing prices, unavailable reviews, malformed product fields, duplicate products, empty review sets, API failures, and model timeouts. The system should never silently substitute fabricated values. If a critical evidence source is unavailable, the final report should explicitly lower the evidence confidence or state that the recommendation cannot be completed reliably.

TABLE VI: Proposed Evaluation Plan

Component	Metrics / evidence	Purpose
Fake-review classifier	Precision, recall, F1, confusion matrix, ROC-AUC where appropriate	Measure credibility classification quality.
Sentiment model	Accuracy/F1 by polarity class	Check opinion classification.
Aspect extraction	Precision, recall, F1 on annotated aspects	Measure feature-level opinion extraction.
Product normalization	Field-level correctness	Ensure comparable product attributes.
Recommendation	Agreement with expert labels / scenario tests	Assess decision consistency.
Explanation	Human rating of clarity and faithfulness	Verify usefulness of reasons.
Latency	Average and percentile response time	Assess practical usability.
Robustness	Missing/noisy/contradictory input tests	Identify failure conditions.

IX. Evaluation Methodology and Validation Plan

Because the supplied BUYSURE AI material defines the intended functionality but does not provide experimental measurements, this paper does not invent accuracy, precision, recall, F1-score, latency, or user-study results. Instead, this section defines a reproducible evaluation plan that can be completed once the implementation and dataset are finalized.

A. Evaluation Dimensions

Evaluation should be divided into four levels. First, component-level evaluation measures the quality of review classification, sentiment analysis, and aspect extraction. Second, data-level evaluation measures whether product information is correctly normalized. Third, decision-level evaluation measures whether recommendations agree with an expert or predefined evaluation protocol. Fourth, user-level evaluation measures whether the report is understandable and useful.

B. Dataset Construction

A credible dataset should contain reviews from multiple products and, ideally, multiple product categories. Labels for suspicious reviews should come from a documented annotation procedure rather than assuming that negative sentiment equals fake content. Review credibility and sentiment are different labels. A review can be negative and genuine, positive and genuine, or suspicious regardless of its sentiment.

For product comparison, a separate evaluation set can contain user scenarios such as budget-only, feature-priority, strict-requirement, and trade-off cases. Human evaluators can determine which product best satisfies each scenario. The system's ranking can then be compared against these reference decisions.

C. Ablation and Comparative Testing

To understand which signals matter, the evaluation should compare configurations. A review-only model can be compared with a review-plus-metadata model. A rating-only baseline can be compared with the complete evidence-fusion system. Similarly, a uniform-weight recommendation can be compared with a personalized-weight recommendation. These comparisons reveal whether the additional system complexity contributes measurable value.

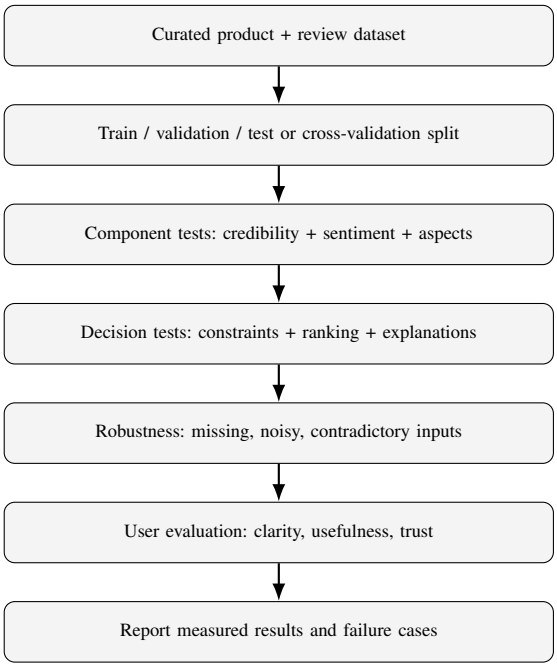


Fig. 5: Proposed validation workflow for the completed BUYSURE AI implementation.

D. Error Analysis

Aggregate accuracy alone is insufficient. The development team should inspect false positives and false negatives in suspicious-review detection. Particular attention should be given to short reviews, multilingual or code-mixed reviews, sarcasm, copied reviews, highly emotional genuine reviews, and AI-generated content. The final report should also be tested for explanation errors, such as citing a negative review that the credibility module has already considered unreliable.

E. Decision Consistency Tests

The recommendation engine should be tested with controlled scenarios. For example, if two products are identical except that one violates a hard budget constraint, the engine should not recommend the over-budget product. If the user's highest priority changes from battery to camera, the ranking should change when the underlying products differ on those attributes. Such tests validate the personalization mechanism independently of model accuracy.

X. Limitations, Ethics, and Risk Management

A. Project-Scope Limitations

The supplied project abstract does not provide a specific dataset, trained-model architecture, numerical benchmark, deployment stack, or completed evaluation. Therefore, the present paper deliberately avoids claiming measured accuracy. The architecture should be read as a detailed engineering design and research methodology for the BUYSURE AI project.

The future features listed in the supplied abstract—image and screenshot analysis, price tracking, seller risk detection, and automatic alternative discovery—are presented as extensions. They should not be described as completed functionality until implemented and evaluated. [1]

B. Ethical Use of Review Detection

Fake-review detection has consequences for trust and reputation. A model prediction should therefore be treated as an evidence signal rather than a final accusation. False positives are especially important because a genuine reviewer may use unusual language, post multiple reviews, or express extreme sentiment. The interface should avoid identifying individual people as fraudulent unless the underlying data and governance process support such a claim.

C. Bias and Domain Shift

A model trained on one category may not transfer equally well to another. Language patterns in electronics reviews may differ from those in clothing, food, appliances, or travel services. Product categories also contain different feature vocabularies. Cross-category validation is therefore important before presenting the system as a general-purpose shopping assistant.

D. Price and Availability Uncertainty

Price is time-dependent. A recommendation that is correct at one time may become less attractive after a price change. If price tracking is not available, the system should label the price as the value observed at analysis time rather than implying a historical trend. Similarly, availability should not be inferred when the system does not have current inventory information.

E. Explainability Requirement

The most important ethical property of the recommendation layer is traceability. A user should be able to identify the major evidence that caused a Buy, Wait, or Avoid result. Explanations should be generated from stored evidence rather than invented after the recommendation is produced. This reduces the risk of plausible-sounding but unsupported explanations.

F. Practical Deployment Considerations

For a deployable version, the system should maintain separate development, evaluation, and production configurations. Model versions, preprocessing rules, feature dictionaries, and scoring weights should be versioned so that a recommendation can be reproduced later. External product data should also carry collection timestamps because price, availability, seller status, and review counts can change independently of the product itself.

The application should expose graceful degradation. If reviews are unavailable, the system can still compare structured product features but should reduce the confidence of review-derived conclusions. If the current price cannot be verified, the report should avoid strong price-timing language. If a model is unavailable, the system should not silently substitute a fabricated result; it should either use a documented fallback or clearly state that the relevant evidence could not be analyzed.

A practical deployment can also maintain an audit record containing the user requirements, products compared, data timestamps, model versions, major evidence signals, and final decision. This record is useful for debugging, evaluation, and future research while still requiring appropriate privacy controls. These deployment principles strengthen the credibility of the system without changing the core BUYSURE AI concept.

XI. Future Enhancements

The project has several clear directions for future development.

A. Product Image and Screenshot Understanding

The supplied abstract proposes image and screenshot analysis. A multimodal model could read product screenshots, identify visible specifications, compare package information, and extract fields that are difficult to obtain from plain text. This module should be treated as another evidence source and should retain confidence and provenance.

B. Real-Time Price Tracking

Price tracking can transform the *Wait* decision from a general caution into a measurable timing recommendation. A future implementation could store price observations, estimate short-term movement, and notify users when a product crosses a target price. Such functionality requires reliable time-stamped price data and careful handling of dynamic pages.

C. Seller Risk Detection

Seller risk can be represented using available seller rating, review patterns, return policy, fulfillment information, warranty conditions, and other permitted marketplace signals. The output should remain a risk indicator rather than a definitive judgment about a seller.

D. Automatic Alternative Discovery

If a selected product fails an important requirement, the system can search the candidate set for products with a better requirement match. An alternative should not be recommended merely because it has a higher numerical score; the report should explain which requirement or trade-off makes the alternative preferable.

E. Multilingual and Code-Mixed Review Analysis

Real-world shopping reviews can contain English, regional languages, abbreviations, transliterated text, and code-mixed sentences. Supporting multilingual and code-mixed review analysis would improve applicability in diverse markets. Such support should be evaluated separately because translation or tokenization errors can alter sentiment and aspect meaning.

F. Continuous Learning and Feedback

Future versions could allow users to indicate whether a recommendation was useful. Such feedback can help improve personalization, but it should not be used to silently change the credibility label of a review. Feedback pipelines require versioning, privacy controls, and monitoring for feedback bias.

XII. Conclusion

BUYSURE AI is proposed as an intelligent shopping assistant that addresses a practical weakness of modern e-commerce: consumers have abundant product information but limited time and limited ability to evaluate its reliability. The supplied project concept identifies fake reviews, misleading information, price variation, unreliable sellers, and large numbers of similar products as key decision difficulties. The proposed architecture responds by combining product-data normalization, NLP-based review analysis, suspicious-review detection, sentiment and aspect extraction, personalized product comparison, evidence fusion, and an explainable Buy/Wait/Avoid decision.

The major design principle is separation of evidence from decision. Reviews are processed into credibility and opinion evidence; product attributes are normalized; user requirements are represented explicitly; and the decision engine combines these signals through a transparent scoring and rule framework. This organization makes the system easier to test, explain, and extend. It also avoids presenting unimplemented future capabilities as current results.

The current project material does not provide measured model performance, so this paper does not invent experimental accuracy or benchmark results. Instead, it provides a detailed validation methodology covering component-level classification, recommendation consistency, robustness, explanation quality, and user usefulness. Once the implementation and dataset are finalized, the proposed evaluation can be used to populate a rigorous results section.

The long-term value of BUYSURE AI lies in converting large volumes of heterogeneous shopping evidence into a compact and understandable decision report. By combining credibility-aware review analysis with personalized product comparison and explicit explanations, the system can serve as a practical decision-support layer between the raw e-commerce marketplace and the final human purchase decision.

For the project team, the next concrete milestone is to connect this architecture to a documented dataset and completed implementation, record the exact algorithms and thresholds used, and execute the evaluation plan described above. The resulting measured evidence can then replace the present proposed-methodology statements with experimentally verified results. This progression keeps the paper academically defensible while leaving a clear path from the current mini-project implementation toward a stronger research prototype.

Acknowledgment

The authors acknowledge the Department of Computer Science and Engineering, Vels Institute of Science, Technology and Advanced Studies (VISTAS), and the project guide for academic guidance and support in developing the BUYSURE AI project concept.

References

[1] S. R. Srikowsalya, S. Sujith, and S. Lingesh, "BUYSURE AI: Project Abstract," Department of Computer Science and Engineering, VISTAS, 2026.

[2] M. A. et al., "Fake Reviews Detection: A Survey," *IEEE Access*, vol. 9, pp. 65771–65802, 2021, doi: 10.1109/ACCESS.2021.3075573.

[3] "Fake Reviews Detection Using Deep Learning: A Survey," in *Proc. Intelligent Methods, Systems, and Applications (IMSA)*, 2024, doi: 10.1109/IMSA61967.2024.10652819.

[4] "Fake Review Detection in E-Commerce Using Machine Learning and NLP Technique," in *Proc. 3rd Int. Conf. Inventive Computing and Informatics (ICICI)*, 2025, doi: 10.1109/ICICI65870.2025.11069636.

[5] "A Hybridized Approach for Enhanced Fake Review Detection," *IEEE Trans. Computational Social Systems*, vol. 11, no. 6, pp. 7448–7466, Dec. 2024, doi: 10.1109/TCSS.2024.3411635.

[6] M. Chelliah and S. Sarkar, "Product Recommendations Enhanced with Reviews," ACM RecSys tutorial material, 2017.

[7] "Multi-Criteria Review-Based Recommender System–The State of the Art," *IEEE Access*, vol. 7, pp. 169446–169468, 2019, doi: 10.1109/ACCESS.2019.2954861.

[8] "A Survey of Explainable E-Commerce Recommender Systems," in *Proc. IEEE/WIC/ACM Int. Joint Conf. Web Intelligence and Intelligent Agent Technology (WI-IAT)*, 2022, doi: 10.1109/WI-IAT55865.2022.00115.

[9] "A Review of Text-Based Recommendation Systems," *IEEE Access*, vol. 9, pp. 31638–31661, 2021, doi: 10.1109/ACCESS.2021.3059312.

[10] "Sentiment Analysis-Based Product Review System for Enhanced Recommendations," in *Proc. Int. Conf. Artificial Intelligence and Data Engineering (AIDE)*, 2025, doi: 10.1109/AIDE64228.2025.10986799.

[11] "A Systematic Study on the Recommender Systems in the E-Commerce," *IEEE Access*, vol. 8, pp. 115694–115716, 2020, doi: 10.1109/ACCESS.2020.3002803.

[12] "Artificial Intelligence in E-commerce Personalization: Algorithms, Customer Analytics, and Intelligent Recommendation Systems," in *Proc. 2nd Int. Conf. Sustainable Computing and Integrated Communication in Changing Landscape of AI (ICSCAI)*, 2026, doi: 10.1109/ICSCAI68849.2026.11649195.

BUYSURE AI Implementation Readiness Checklist

- Finalize the product/review dataset and document its source and collection date.
- Record the exact NLP and ML algorithms, preprocessing rules, and model versions.
- Verify review-credibility labels with an explicit annotation procedure.
- Document all product-feature normalization rules and user-preference weights.
- Implement and test the Buy/Wait/Avoid decision thresholds against controlled scenarios.
- Add explanation traceability so every major recommendation reason maps to stored evidence.
- Evaluate component accuracy, robustness, latency, recommendation consistency, and user clarity.
- Replace proposed-methodology statements with measured results before conference submission.